

# *Hash Join Algorithms*

Instructor: Sharma Chakravarthy


[sharma@cse.uta.edu](mailto:sharma@cse.uta.edu)

The University of Texas @ Arlington

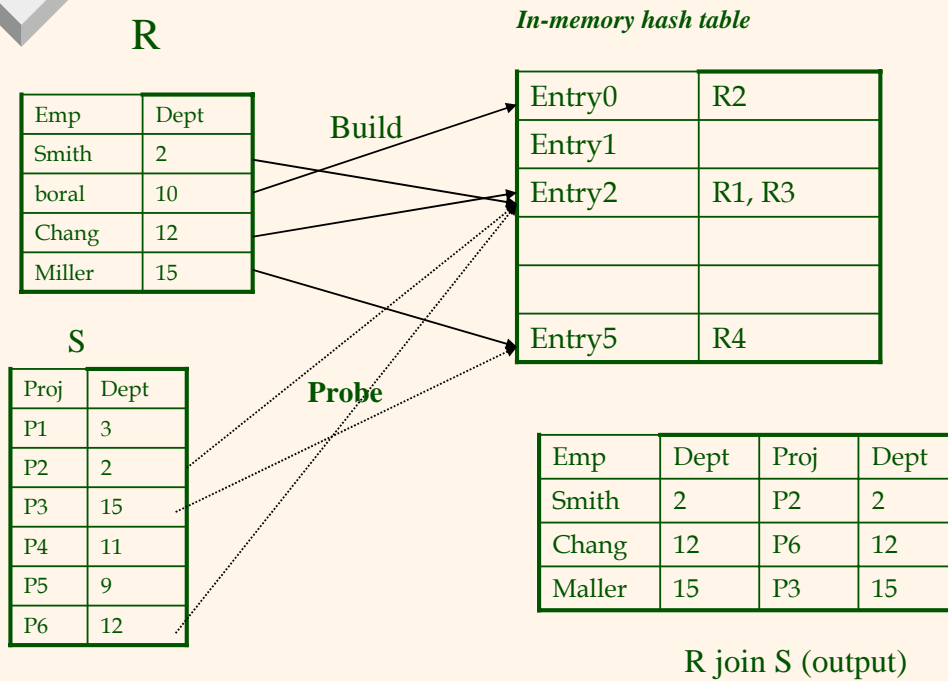


## *Hash-Join Algorithms*

- ❖ In-memory Hash join
  - When you can hold one of the 2 relations in memory
- ❖ Simple hash-based join
  - Efficient when memory is large
  - Too many I/O operations when memory is small
- ❖ GRACE hash-based join
  - Separate partitioning and join phases
  - Easy to **parallelize**
  - Avoids bucket overflow
- ❖ Hybrid hash-based join
  - Combines Basic and Grace hash-join
  - Better memory usage



# In-memory hash-join Algorithm






## Complexity

- ❖ Build phase
  - Read R once and construct in-memory hash table
  - I/Os: M (# of pages of R)
- ❖ Probe phase
  - Read all of S and search for matching tuples
  - I/Os: N (# of pages of S)
- ❖ Total Cost:  $O(M+N)$  if we have enough memory to hold one relation in memory
- ❖ How do you choose the relation for Build?
- ❖ How do you choose the relation for probe?
- ❖ What if we do not have enough memory?



## *Simple Hash Join algorithm*

- ❖ Use whatever memory is available as buckets of **one in-memory hash table** and write the rest to disk
- ❖ Repeat this process until the entire join is performed
- ❖ **Disadvantages:** introduces too many I/O operations when the memory is not too large!
- ❖ Cost:  $O(b \cdot (M+N))$  where  $b$  is the number of **buckets (range of hash function)**!




## *Simple Hash Join Algorithm*

```
/* h is the hash function; h[0..n] is the range of hash function */
/* R[0..n] and S[0..n] are buckets */
i=0; do
for ( each tuple r in R){
    if (h(r) in current_range)
        insert r into the in-memory hash table;
    else write r into R_temp;}
for (each tuple s in S){
    if (h(s) is in current_range{
        use s to probe the in-memory hash table;
        If (any match is found) output the matching tuples;
    else write s into S_temp; }
    R = R_temp;
    S = S_temp;
    current_range = h[i+1];
}
While (R_temp is not empty and S_temp is not empty);
```



## Complexity

- ❖ Let size of R be M pages; size of S be N pages
- ❖ Let the hash function divide them uniformly into b buckets
- ❖ If you have b hash buckets for the simple hash join algorithms, then you need  $b * (M+N)$  I/O's (Try to derive this expression!)
- ❖ You read and write each relation b times!
- ❖ Typically, b ranges from 10 to 1024 or even larger
- ❖ How can we reduce it further?
- ❖ How many buffer pages do we need



## *GRACE Hash Join Algorithm*

### Partitioning phase

- ❖ Apply a hash function  $h(x)$  to the join attributes of the tuples in **both** R and S. Assume  $b$  buckets
- ❖ According to the hash value, each tuple is put into a corresponding bucket. Write these buckets to disk as separate files.


### Joining phase:

- ❖ Use the basic hash-join algorithm
- ❖ Get one partition of R and the **corresponding** partition of S and apply the basic hash algorithm using a **different hash function**. Why?

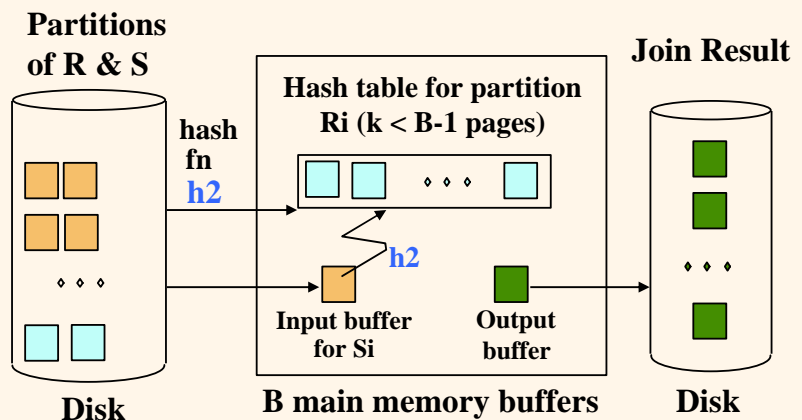


# Hash-Join

- ❖ Partition both relations using hash fn **h**: R tuples in partition *i* will only match S tuples in partition *i*.



- ❖ Read in a partition of R, hash it using **h2** ( $\neq h$ !). Scan matching partition of S, search for matches.





## *Grace Hash Join*

- ❖ Range of  $H(x)$  is  $1, \dots, N$
- ❖  $R_1, \dots, R_n$  and  $S_1, \dots, S_n$  are **disjoint** subsets of  $R$  and  $S$
- ❖  $R$  is the Union  $(R_1, \dots, R_n)$  and  $S$  is the union  $(S_1, \dots, S_n)$
- ❖ We need to join **only**  $R_i$  with  $S_i$ . Why?
- ❖ The efficiency comes from the reduction in work load which is illustrated below.



## *Grace Hash Join Algorithm*


```
/* h[1..n]: range of hash function; R[1..n] and S[1..n] are buckets */
for ( each tuple r in R){
    apply hash function to the join attributes of r;
    put r into the appropriate bucket R[i]}

for (each tuple s in S){
    apply hash function to the join attributes of s;
    put r into the appropriate bucket S[i]}


for (i=1; i <= n; i++){
    build the hash table for R[i]; /* using a different hash function h2*/
    for (each tuple s in S[i]){
        apply the hash function h2 to the join attributes of S;
        use s to probe the hash table;
        output any matches to the result relation;}
    }
```



## *Workload in hash join*



Nested loop join



Grace hash join



## *Observations on Hash-Join*

- ❖ Given  $B$  buffer pages, the maximum # of partitions is  $B-1$
- ❖ Assuming that partitions are of equal size, the size of each  $R$  partition is  $M/(B-1)$
- ❖ The number of pages in the (in-memory) hash table built during the building phase is  $f \cdot M/(B-1)$  where  $f$  is the fudge factor
- ❖ During the probing phase, in addition to the hash table for the  $R$  partition, we require a buffer page for scanning the  $S$  partition, and an output buffer.
- ❖ Therefore, we require  $B > f \cdot M/(B-1) + 2$
- ❖ Approximately, we need  $B > \sqrt{M}$  for the hash join algorithm to perform well.



## *Observations on Hash-Join*


If we build an in-memory hash table to speed up the matching of tuples, a little more memory is needed.

- ❖ If the hash function does not partition uniformly, one or more R partitions may not fit in memory. Can apply hash-join technique recursively to do the join of this R-partition with corresponding S-partition.



## *Cost of Grace Hash Join*

- ❖ In partitioning phase,
  - read+write both relations; that is,  $2(M+N)$ .
  - In matching phase, read both relations; that is,  $M+N$  I/Os.
  
- ❖ In our running example, this is a total of 4500 I/Os.



## *Sort-merge join vs. Hash Join*

- ❖ If partitions in hash join are not uniformly sized, hash join could cost more
- ❖ If the available number of buffers falls between  $\sqrt{M}$  and  $\sqrt{N}$ , hash join costs less than sort-merge, since we need enough memory to hold partitions of the smaller relation. Sort-merge buffer needs are based on the larger relation.
- ❖ Hash Join is superior on this count if relation sizes differ greatly. Also, Hash Join shown to be highly parallelizable.
- ❖ Sort-Merge less sensitive to data skew; result is sorted.





## General Join Conditions

- ❖ Equalities over several attributes (e.g., *R.sid=S.sid* AND *R.rname=S.sname*):
  - For Index NL, build index on *<sid, sname>* (if S is inner); or use existing indexes on *sid* or *sname*.
  - For Sort-Merge and Hash Join, sort/partition on combination of the two join columns.
- ❖ Inequality conditions (e.g., *R.rname < S.sname*):
  - For Index NL, need (clustered!) B+ tree index.
    - ◆ Range probes on inner; # matches likely to be much higher than for equality joins.
  - Hash Join, Sort Merge Join not applicable.
  - Block NL quite likely to be the best join method here.



## Hybrid Hash Join Algorithm

```
/* H[0..n] is the range of hash; R[0..n] and S[0..n] are buckets */
for ( each tuple in R){
    if (hash value of r is in H[0])
        insert r into the in-memory hash table;
    else put r into the appropriate bucket R[i];}
for (each tuple s in S){
    if (hash value of s is in H[0]){
        use s to probe the hash table;
        put any matching tuples into the result relation;}
    else put s into appropriate bucket S[i];}
for (i=1; i<=n; i++){
    build the hash table from R[i];
    for (each tuple s in S[i]){
        apply hash function to the join attributes of s;
        use s to probe the hash table;
        output any matches to the result relation;}
}
```



## *Pointer Based Joins*

1. Links represent a limited form of pre-computed results (OO has rekindled this concept)
2. Modeled as TID joins in Ingres
3. Shekita and Carey experiment – 3 pointer based join methods: Nested Loops, Merge-Join and Hybrid-Hash Join
  1. Tuples of R has a pointer to an embedded S tuple
    - Scan R and retrieve S
    - Sort R on the pointers (according to the disk address they point to) and then retrieve all S items in one elevator pass over the disk, reading S page at most once



## *Pointer based Joins(contd)*

- Hybrid-hash join: Partitions relation R on the pointer values ensuring that R tuples with S pointers to the same page are brought together, and then retrieve S pages and tuples
- ❖ Direction of pointers fix the role of relations! (usually, the smaller relation is used for the build phase)
- ❖ Maintenance effort is to be taken into account as well.

## Alternative Join methods

- ❖ S is 10 times R, Memory size 100Kb
- ❖ Cluster Size is 8Kb, Merge fan-in and partitioning fan-out are 10, # of R records/cluster is 20

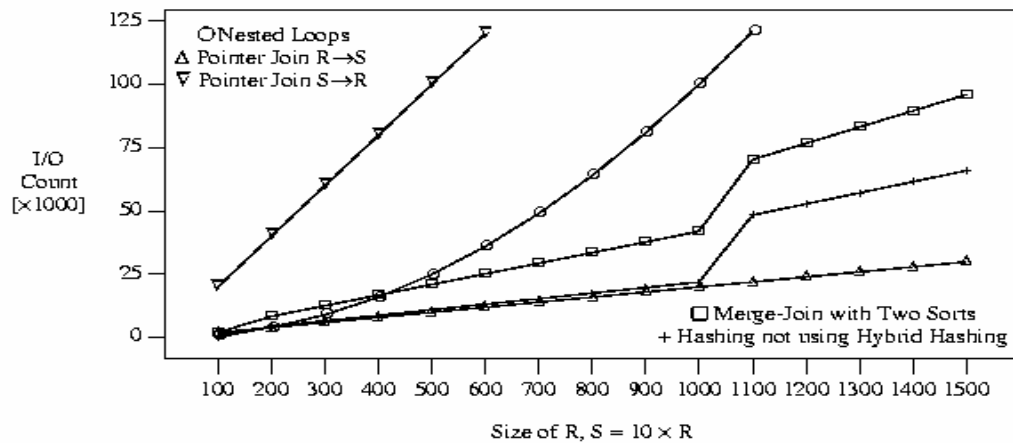



Figure 16. Performance of Alternative Join Methods.



## *Conclusions*

- ❖ Nested Loop joins are unsuitable for medium size and large relations
- ❖ sort based join is not as fast as hash join (merge levels are determined individually for each file, but only the smaller relation determines partition depth)
- ❖ The step is because additional partitioning or merge levels become necessary at that point



## *Aggregation and Duplicate Removal*

- ❖ Surprisingly, a lot in common
- ❖ In one, duplicates are discarded whereas in the other, some computation (e.g., COUNT, SUM, AVG) is performed before discarding the tuple



## *Aggregation and Duplicate Removal*

- ❖ Scalar aggregates compute a single scalar value; from a unary input relation (count of all employees)
  - requires only one pass over data set
  - indices can be exploited where possible (for max, min, count)
- ❖ Aggregate functions determine a set of values from a binary input relation; e.g., sum of salaries for each department
- ❖ The result is a relation (closure property)





## *“Duality” of Sorting and Hashing*

- ❖ Both do approx the same amount of I/O
- ❖ Mirror-images in terms of sequentiality of phase 2
- ❖ Sort-based algorithms
  - Large data sets are divided into subsets using physical rule (into chunks as large as memory)
- ❖ Hash-based algorithms
  - Large data sets are divided into subsets using a logical rule (hash values)
- ❖ Handling large inputs
  - Multi-pass sort vs. recursive partitioning hash
- ❖ It actually goes deeper than this